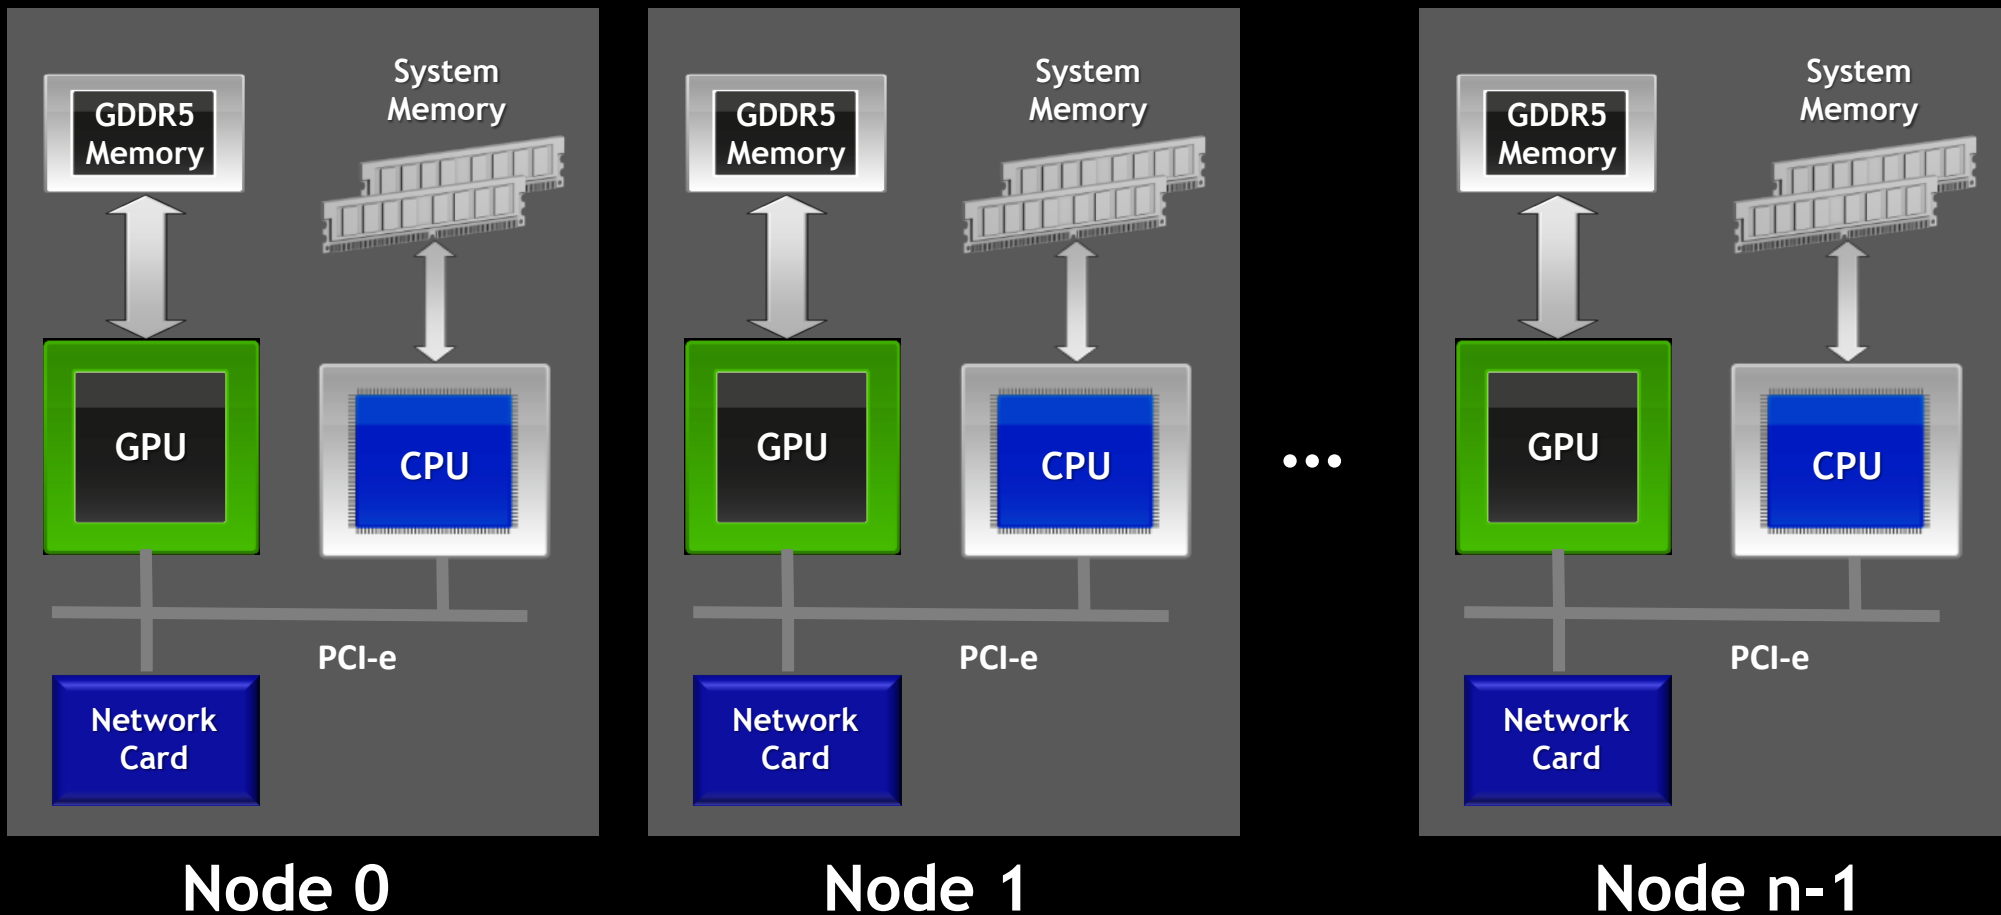


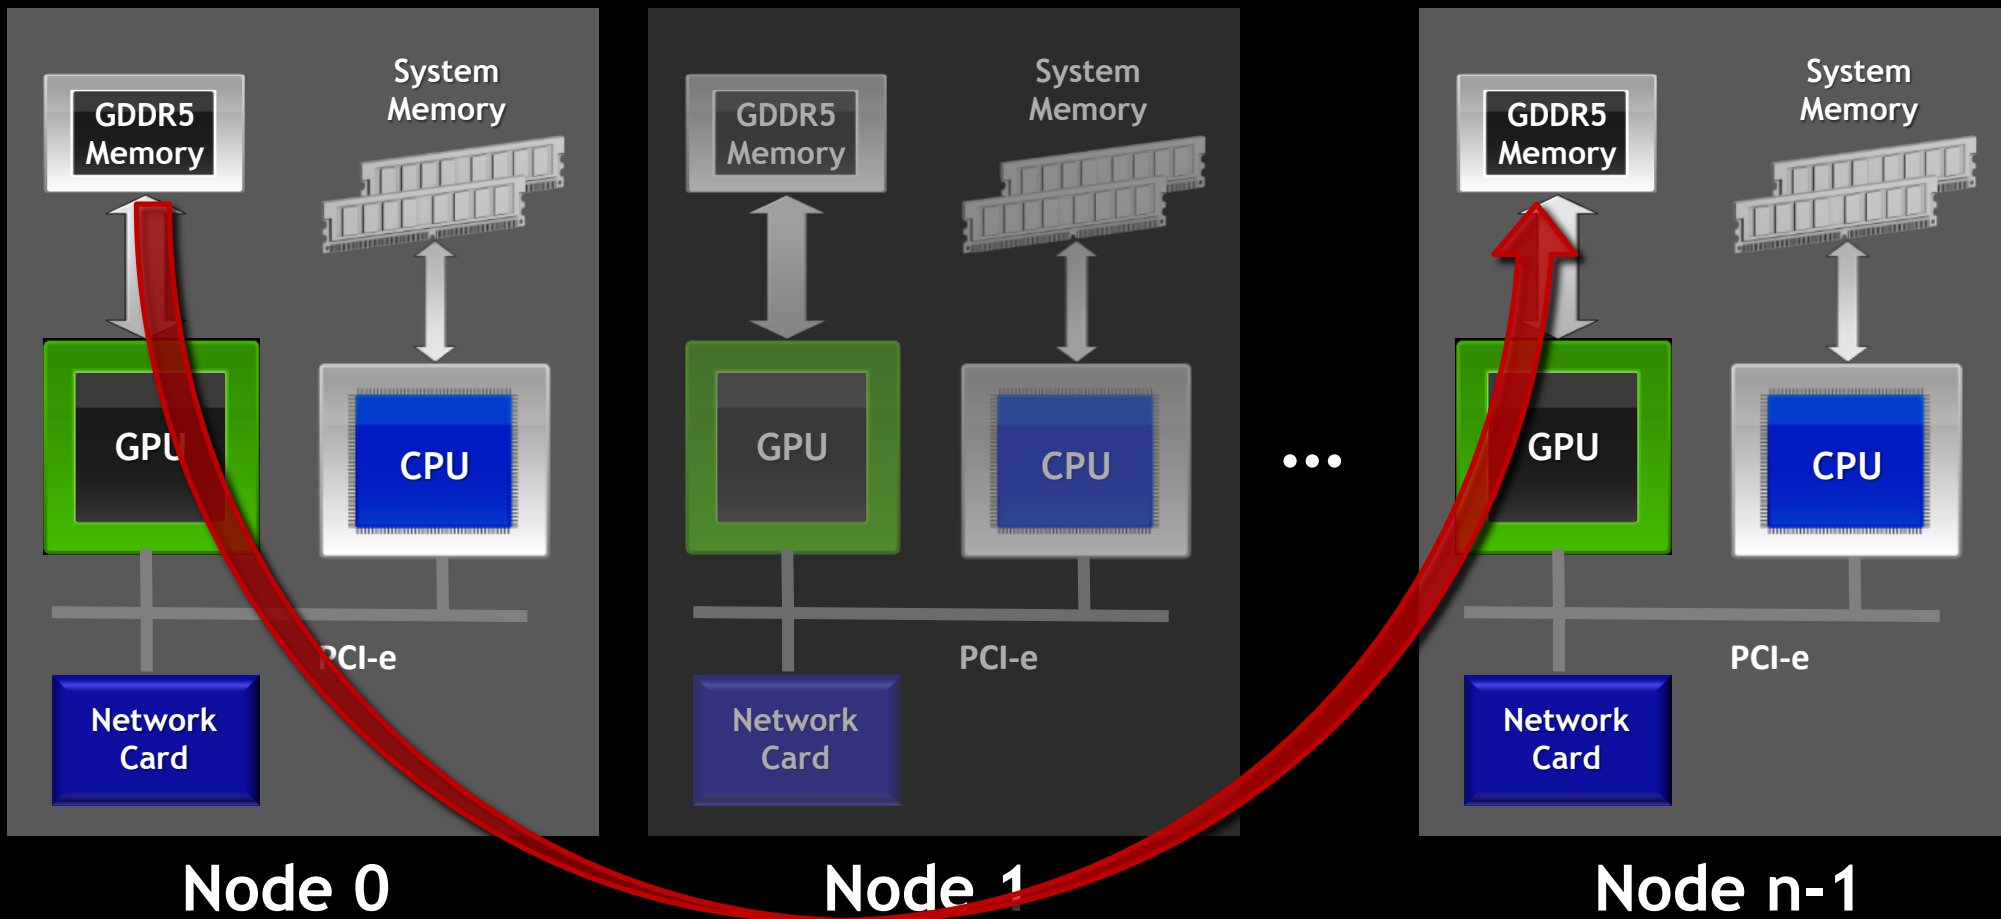
Introduction to CUDA-aware MPI and NVIDIA GPUDirect™

Jiri Kraus
Developer Technology HPC, NVIDIA


MPI+CUDA



MPI+CUDA



MPI+CUDA



```
//MPI rank 0
```

```
MPI_Send(s_buf_d,size,MPI_CHAR,n-1,tag,MPI_COMM_WORLD);
```

```
//MPI rank n-1
```

```
MPI_Recv(r_buf_d,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
```

Outline

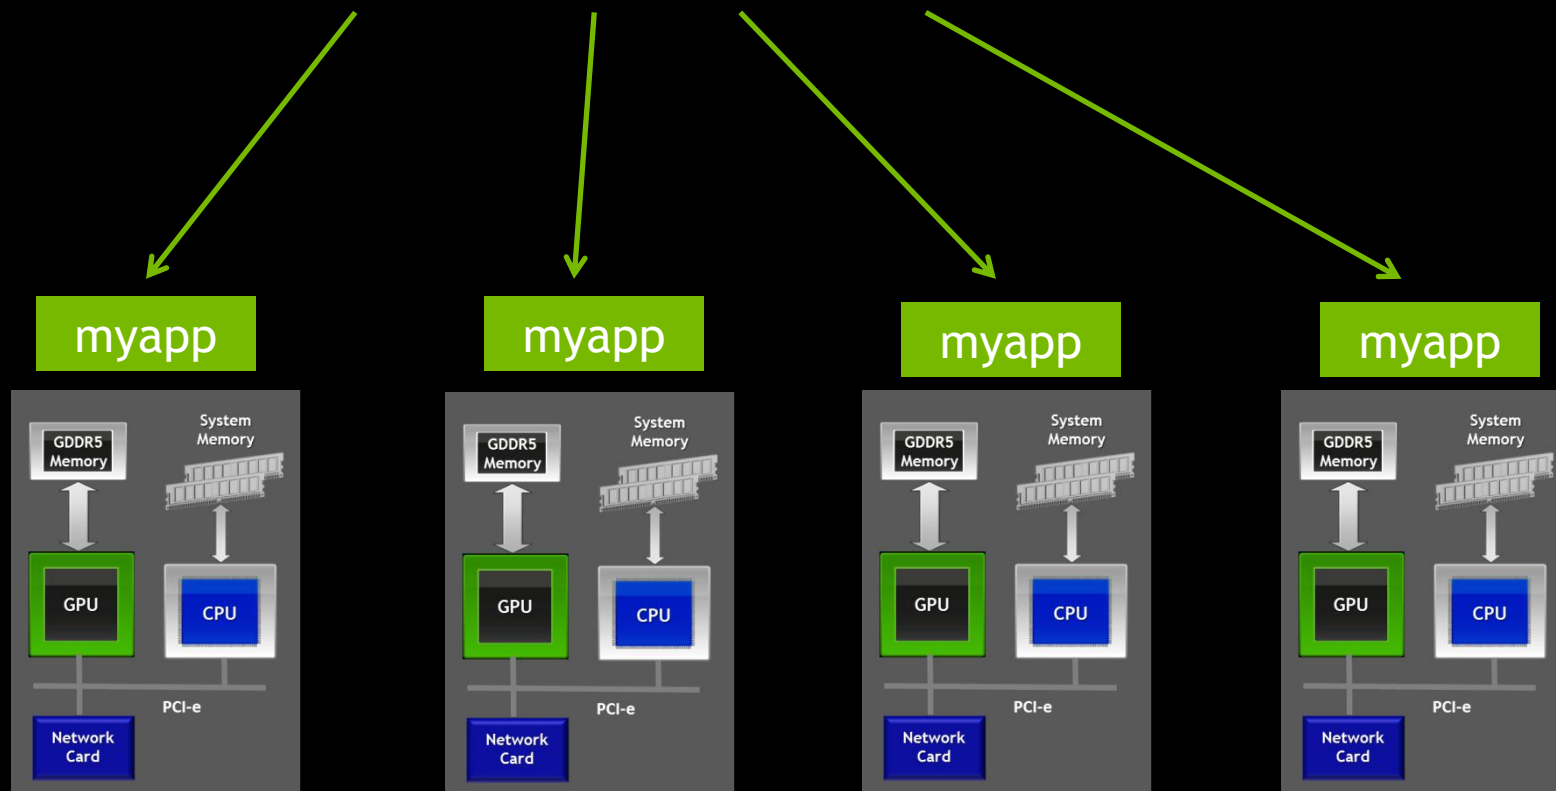
- Short Introduction to MPI
- Unified Virtual Addressing and GPUDirect
- How CUDA-aware MPI works
- Performance Results
- Wrap-up and conclusions

Message Passing Interface - MPI

- Standard to exchange data between processes via messages
 - Defines API to exchanges messages
 - Pt. 2 Pt.: e.g. MPI_Send, MPI_Recv
 - Collectives, e.g. MPI_Reduce
- Multiple implementations (open source and commercial)
 - Binding for C/C++, Fortran, Python, ...

MPI - How to launch a MPI program

```
mpirun -np 4 ./myapp <args>
```




MPI - A minimal program

```
#include <mpi.h>


int main(int argc, char *argv[]) {
    int myrank;
    /* Initialize the MPI library */
    MPI_Init(&argc,&argv);
    /* Determine the calling process rank */
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    /* Call MPI routines like MPI_Send, MPI_Recv, ... */
    /* Shutdown MPI library */
    MPI_Finalize();
    return 0;
}
```


Unified Virtual Addressing

No UVA: Multiple Memory Spaces




UVA : Single Address Space




Unified Virtual Addressing

No UVA: Multiple Memory Spaces



UVA : Single Address Space



- One address space for all CPU and GPU memory
 - Determine physical memory location from a pointer value
 - Enable libraries to simplify their interfaces (e.g. MPI and cudaMemcpy)
- Supported on devices with compute capability 2.0 for
 - 64-bit applications on Linux and on Windows also TCC mode


MPI+CUDA

With UVA and CUDA-aware MPI

```
//MPI rank 0  
MPI_Send(s_buf_d,size,...);
```

```
//MPI rank n-1  
MPI_Recv(r_buf_d,size,...);
```

CUDA-aware MPI makes MPI+CUDA easier.



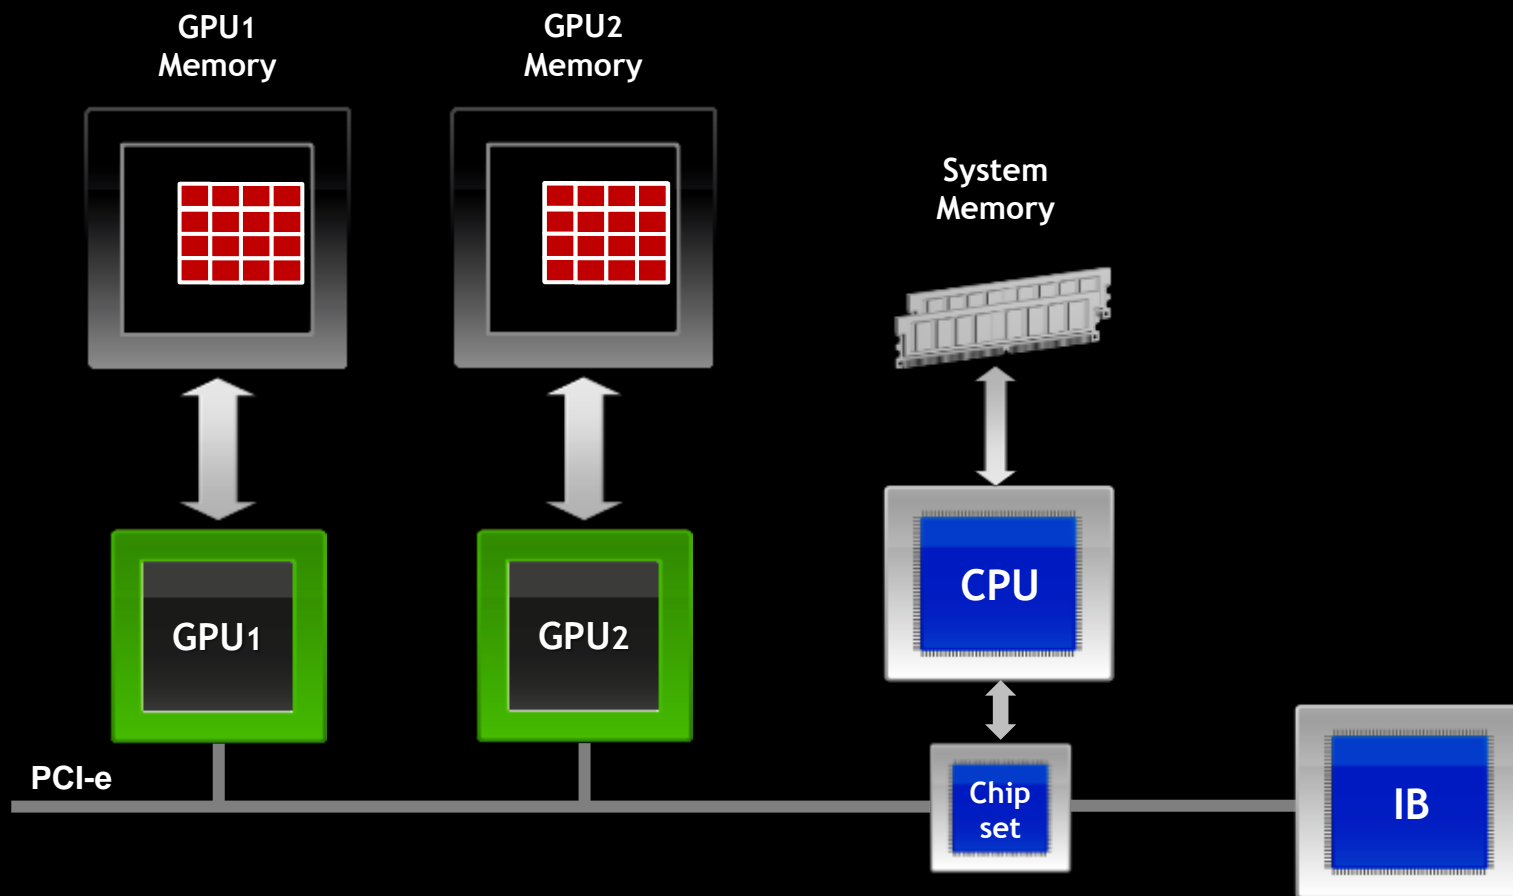
No UVA and regular MPI

```
//MPI rank 0  
cudaMemcpy(s_buf_h,s_buf_d,size,...);  
MPI_Send(s_buf_h,size,...);
```

```
//MPI rank n-1  
MPI_Recv(r_buf_h,size,...);  
cudaMemcpy(r_buf_d,r_buf_h,size,...);
```

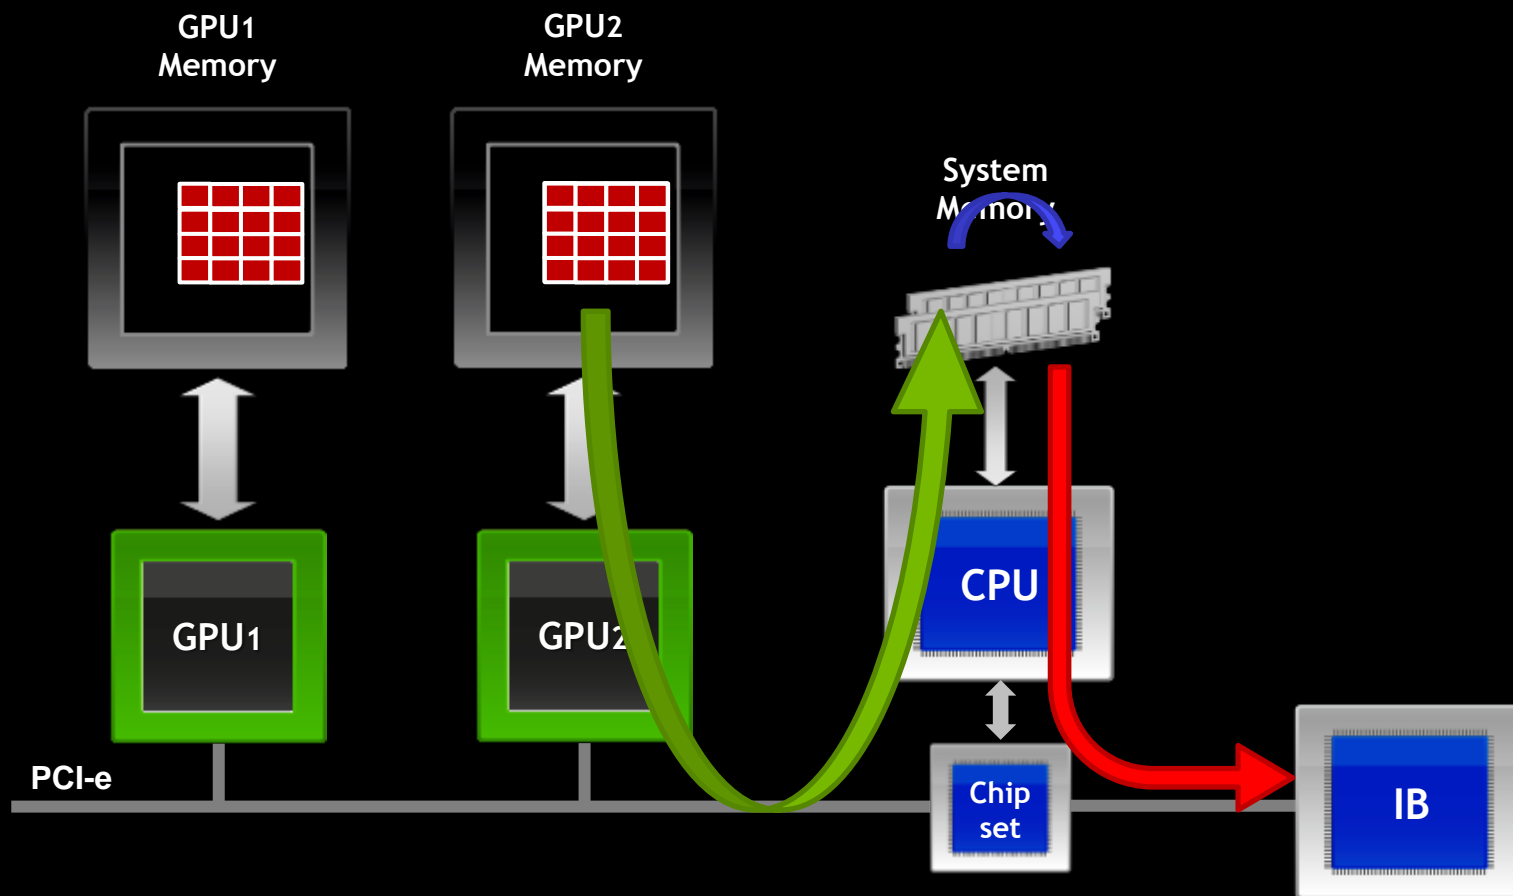
NVIDIA GPUDirect™

Accelerated communication with network & storage devices



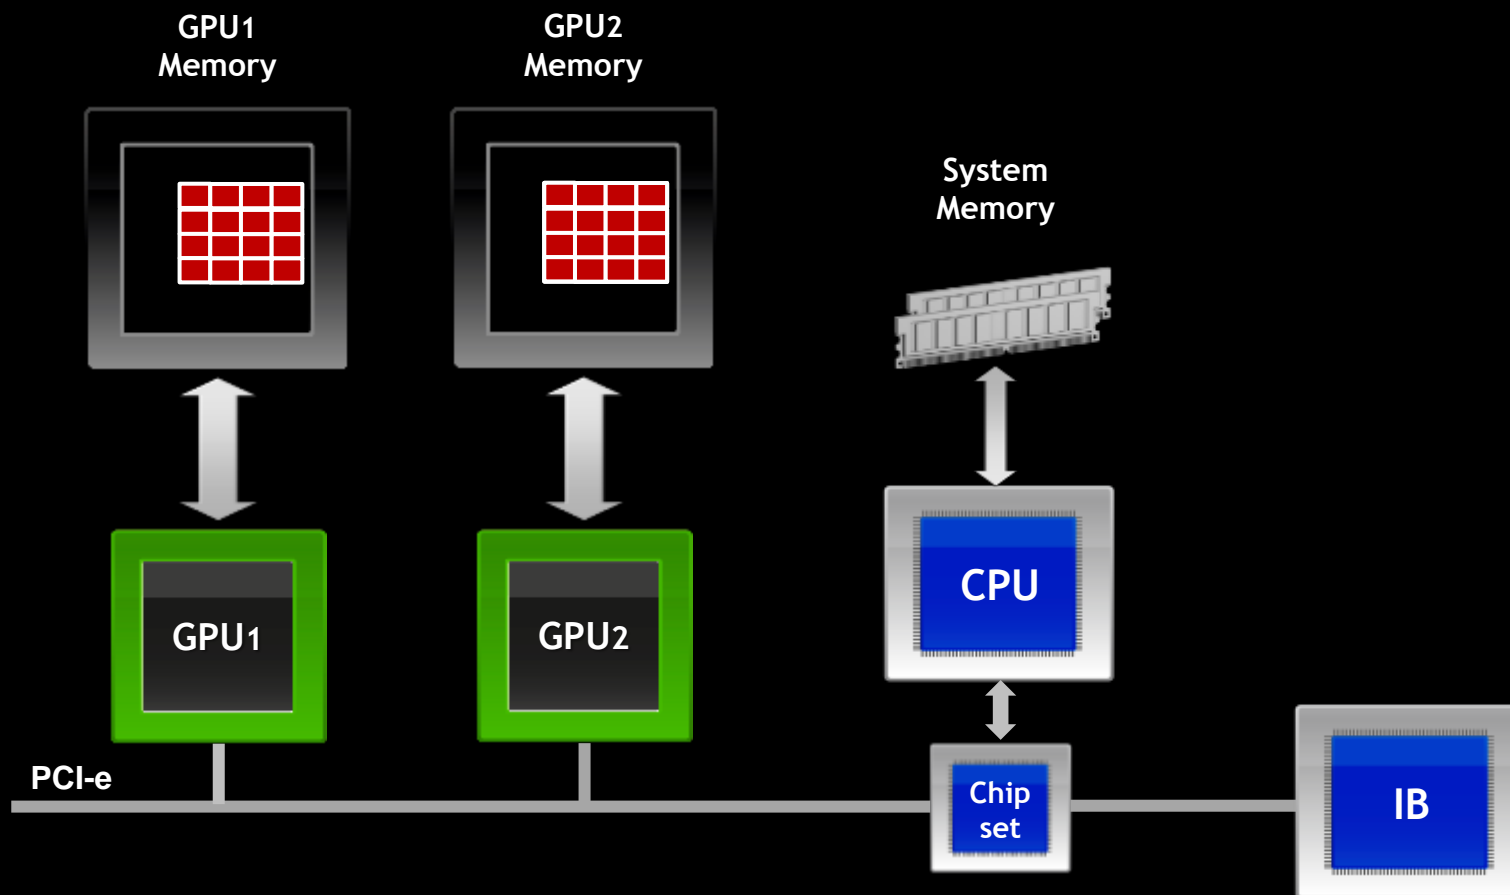
NVIDIA GPUDirect™

Accelerated communication with network & storage devices



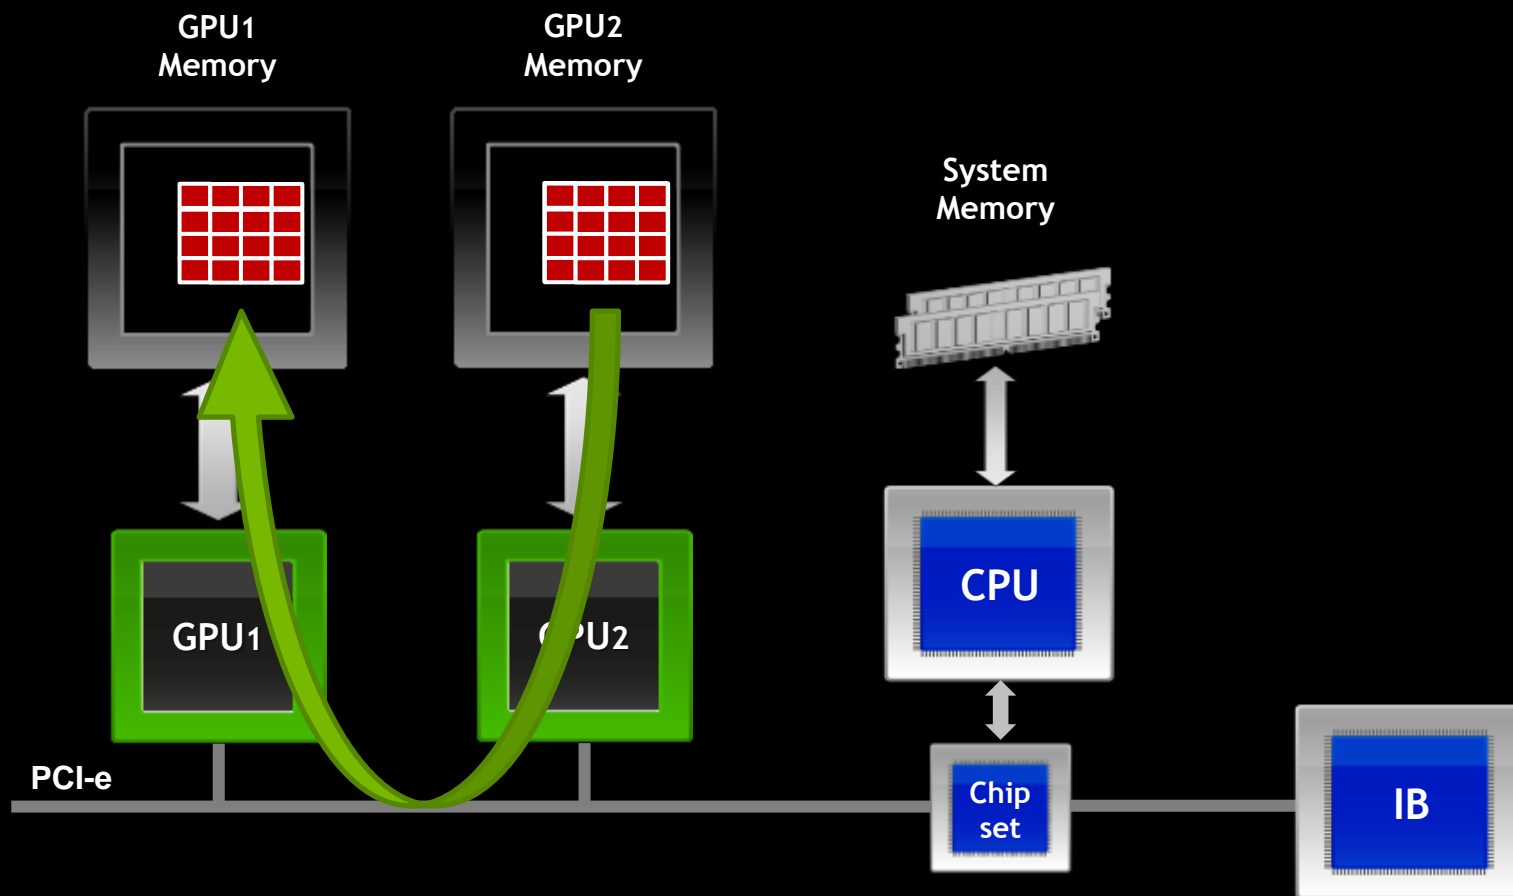
NVIDIA GPUDirect™

Peer to Peer Transfers



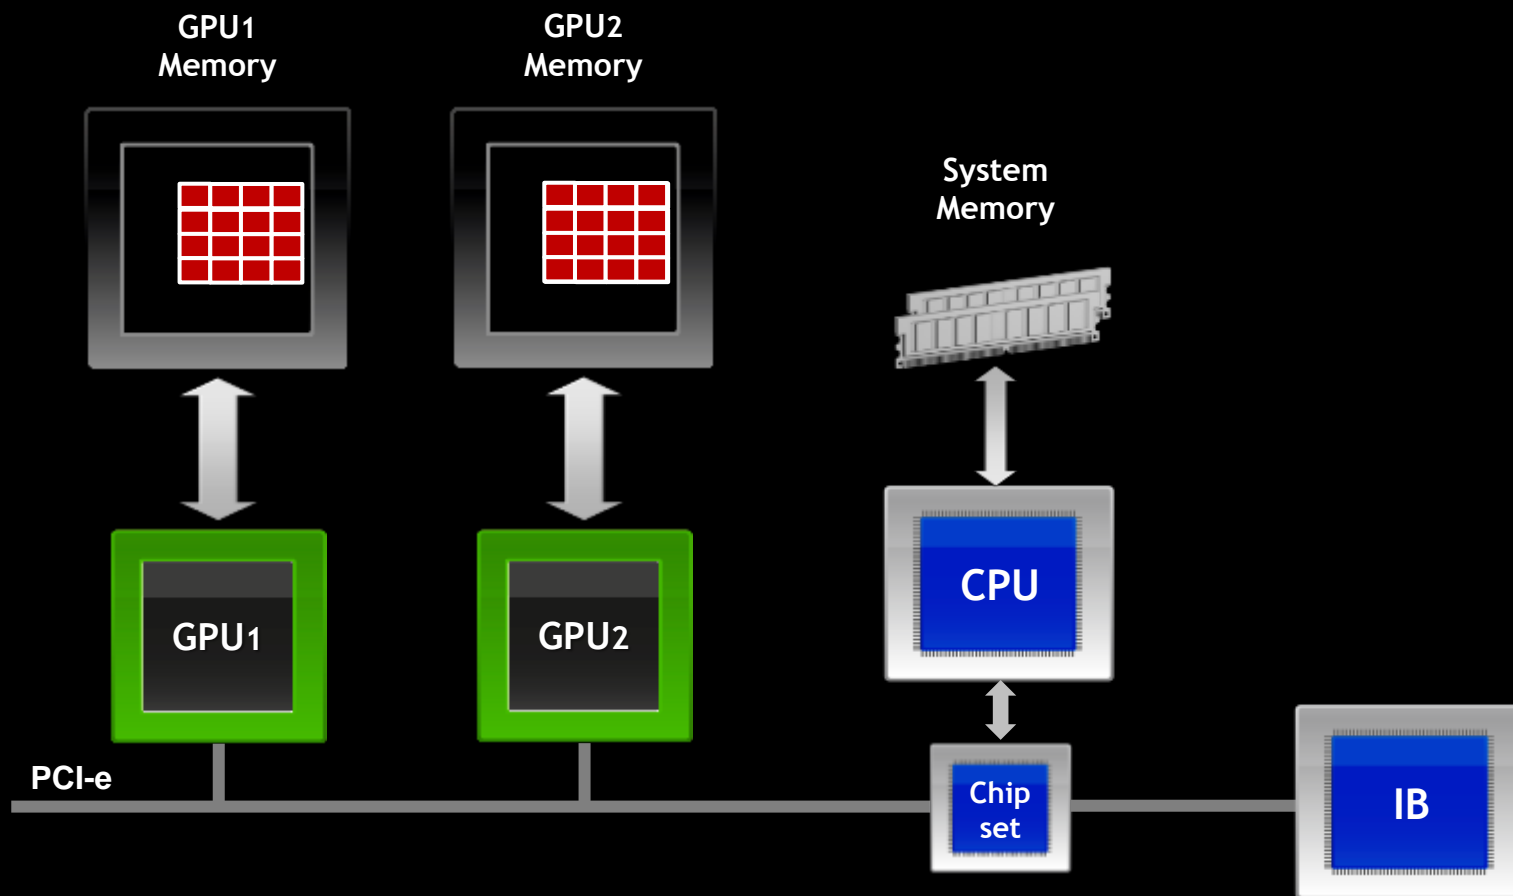
NVIDIA GPUDirect™

Peer to Peer Transfers



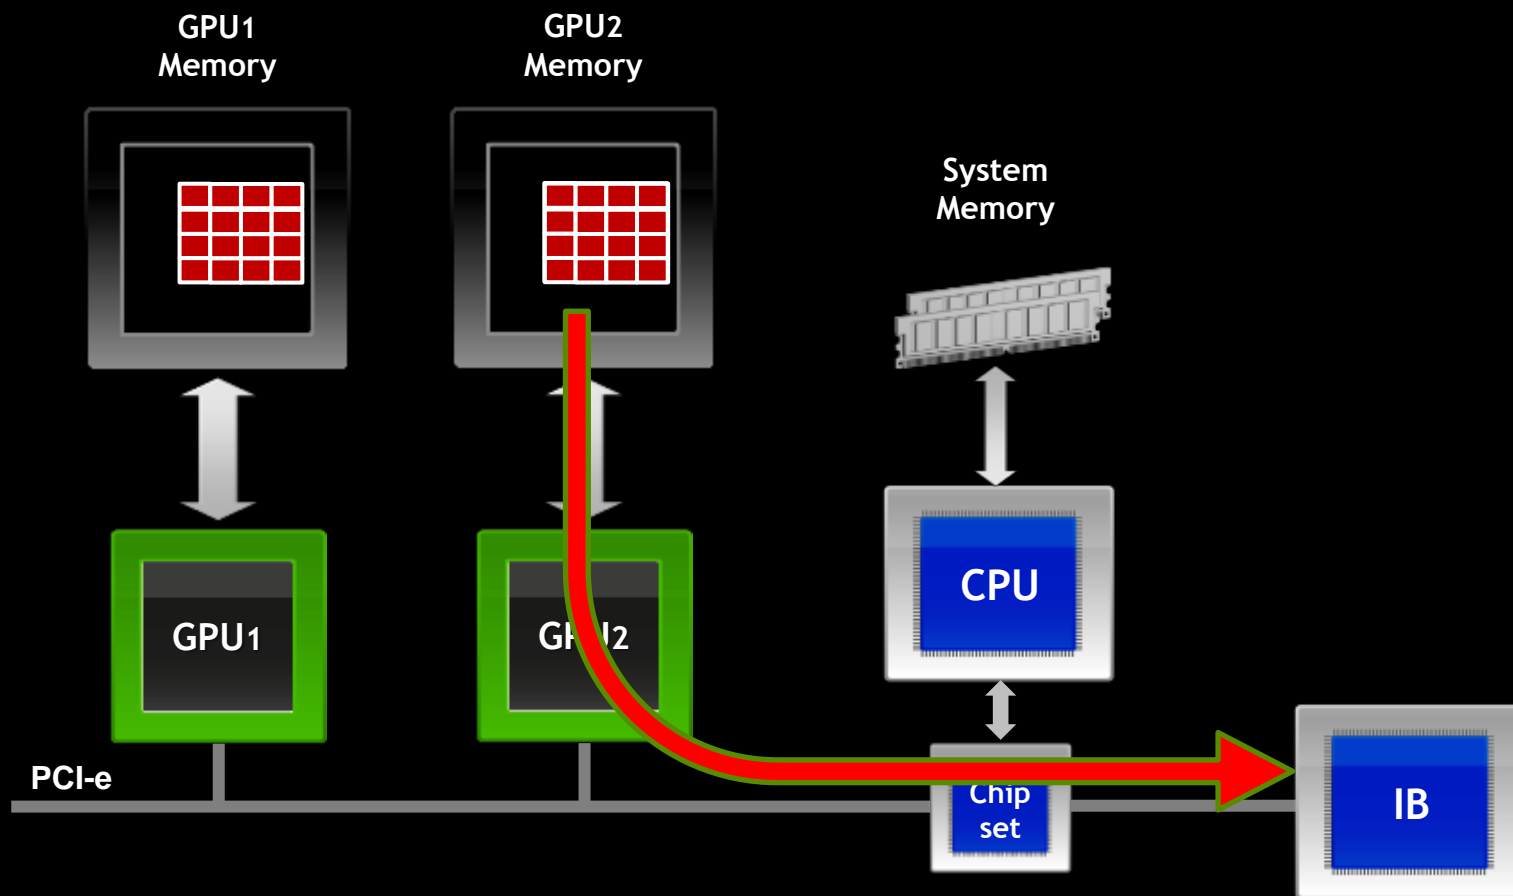
NVIDIA GPUDirect™

Support for RDMA



NVIDIA GPUDirect™

Support for RDMA



CUDA-Aware MPI

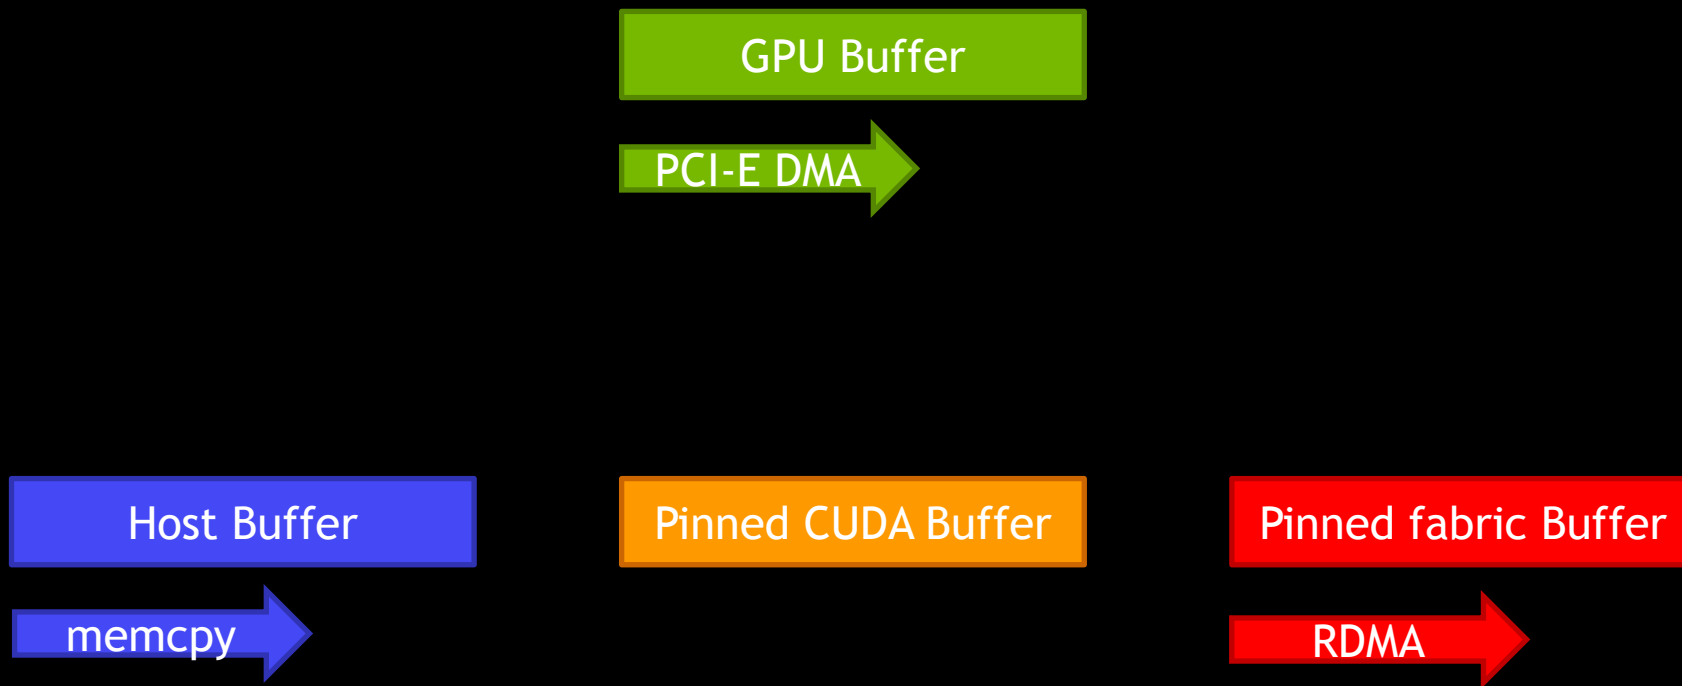
Example:

MPI Rank 0 MPI_Send from GPU Buffer

MPI Rank 1 MPI_Recv to GPU Buffer


- Show how CUDA+MPI works in principle
 - Depending on the MPI implementation, message size, system setup, ... situation might be different
- Two GPUs in two nodes

CUDA-Aware MPI



MPI GPU to Remote GPU

GPUDirect Support for RDMA




```
MPI_Send(s_buf_d,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
```


```
MPI_Recv(r_buf_d,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
```


MPI GPU to Remote GPU

GPUDirect Support for RDMA



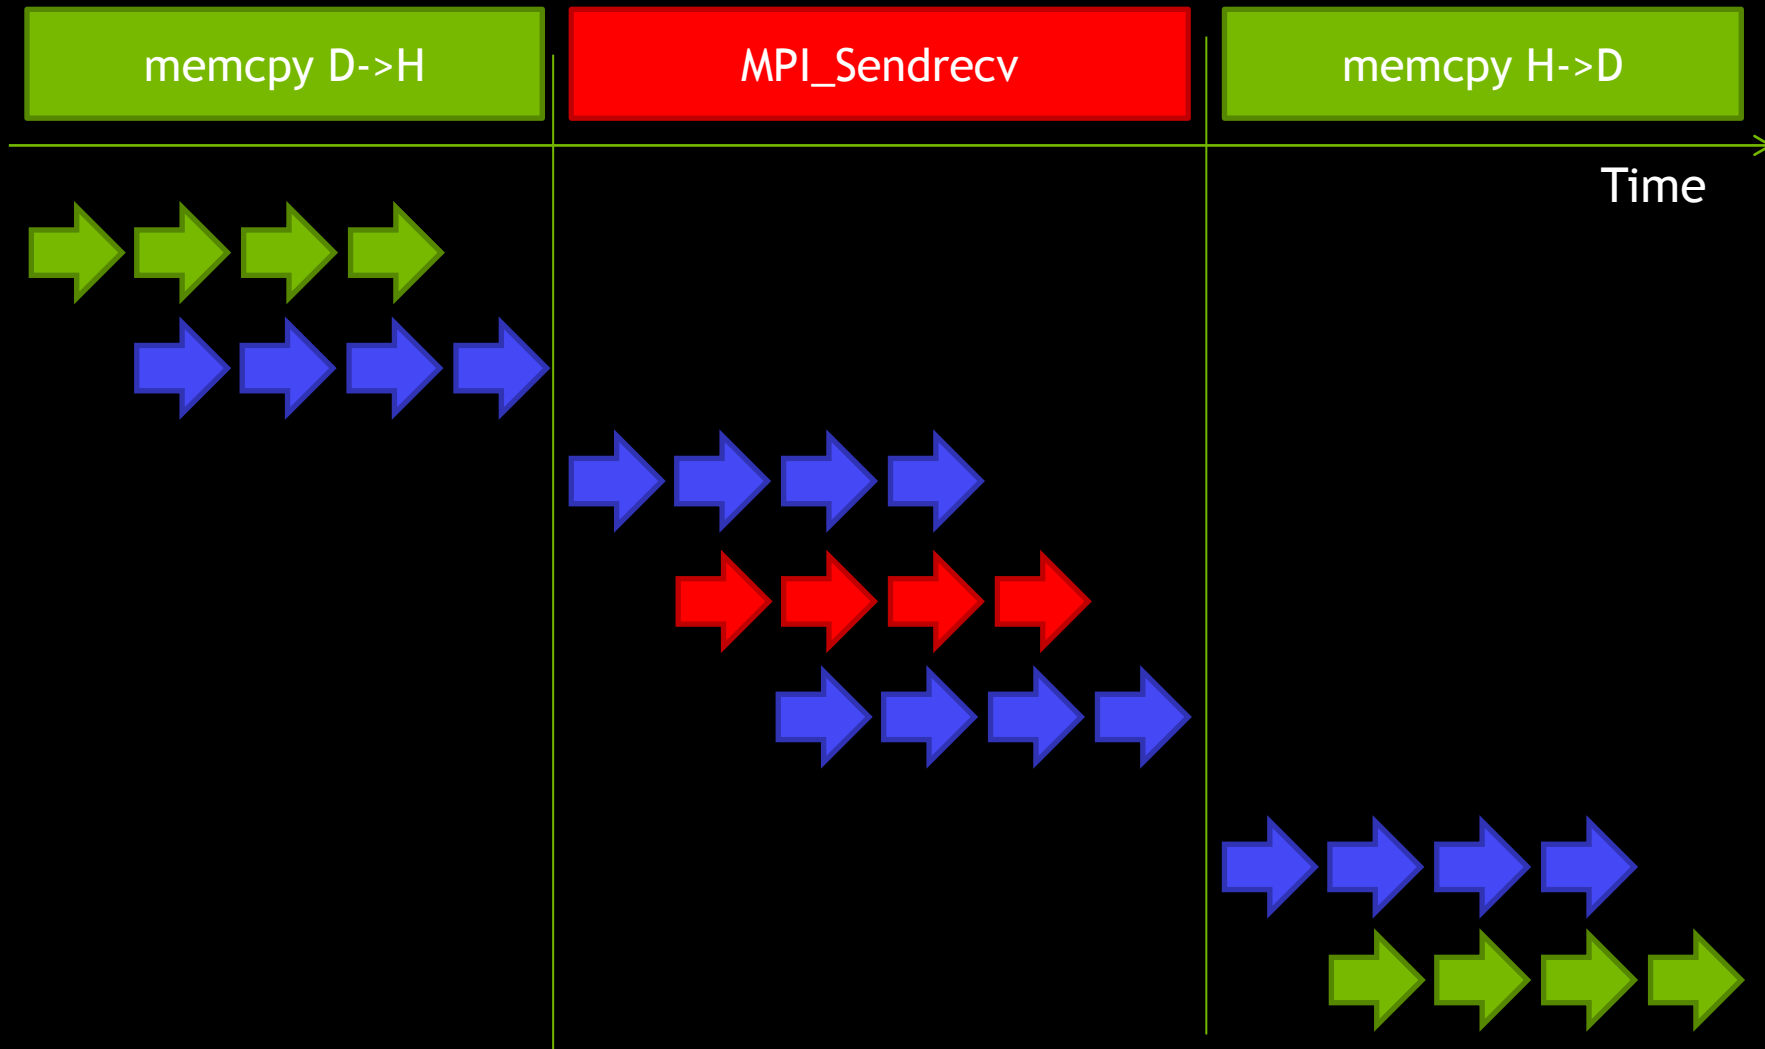
Regular MPI GPU to Remote GPU



```
cudaMemcpy(s_buf_h,s_buf_d,size,cudaMemcpyDeviceToHost);  
MPI_Send(s_buf_h,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
```


```
MPI_Recv(r_buf_h,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);  
cudaMemcpy(r_buf_d,r_buf_h,size,cudaMemcpyHostToDevice);
```

Regular MPI GPU to Remote GPU



MPI GPU to Remote GPU

Without GPUDirect

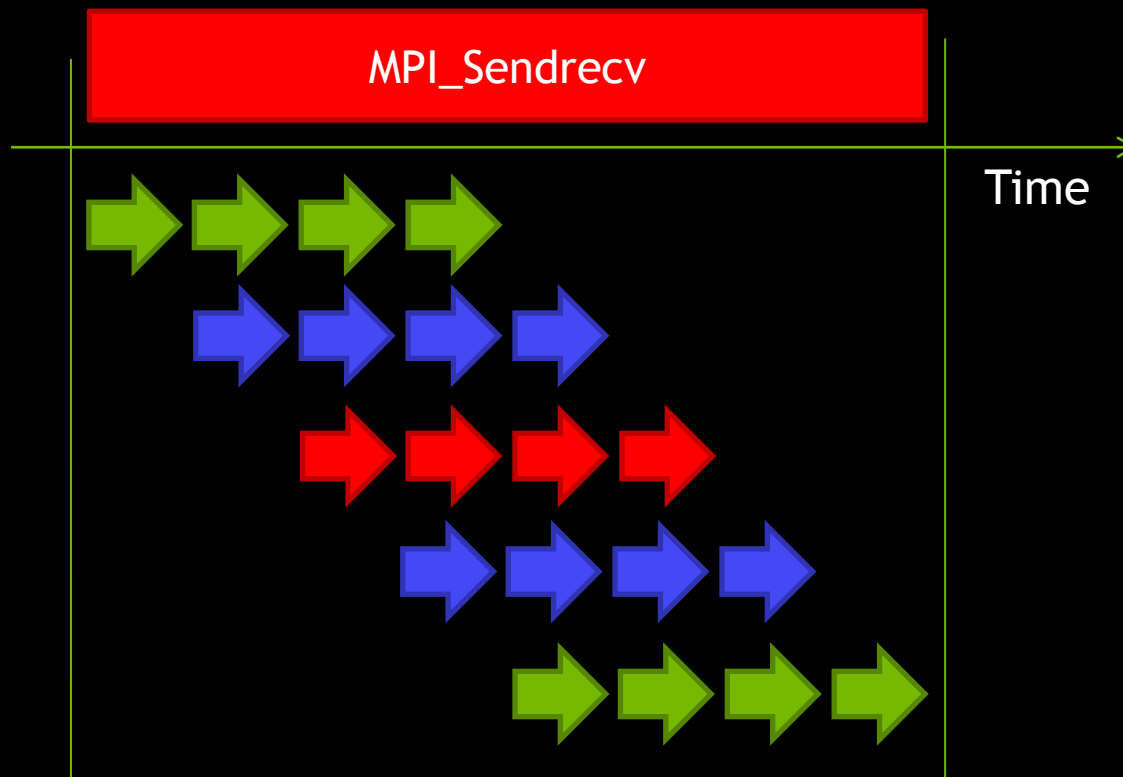


```
MPI_Send(s_buf_h,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
```

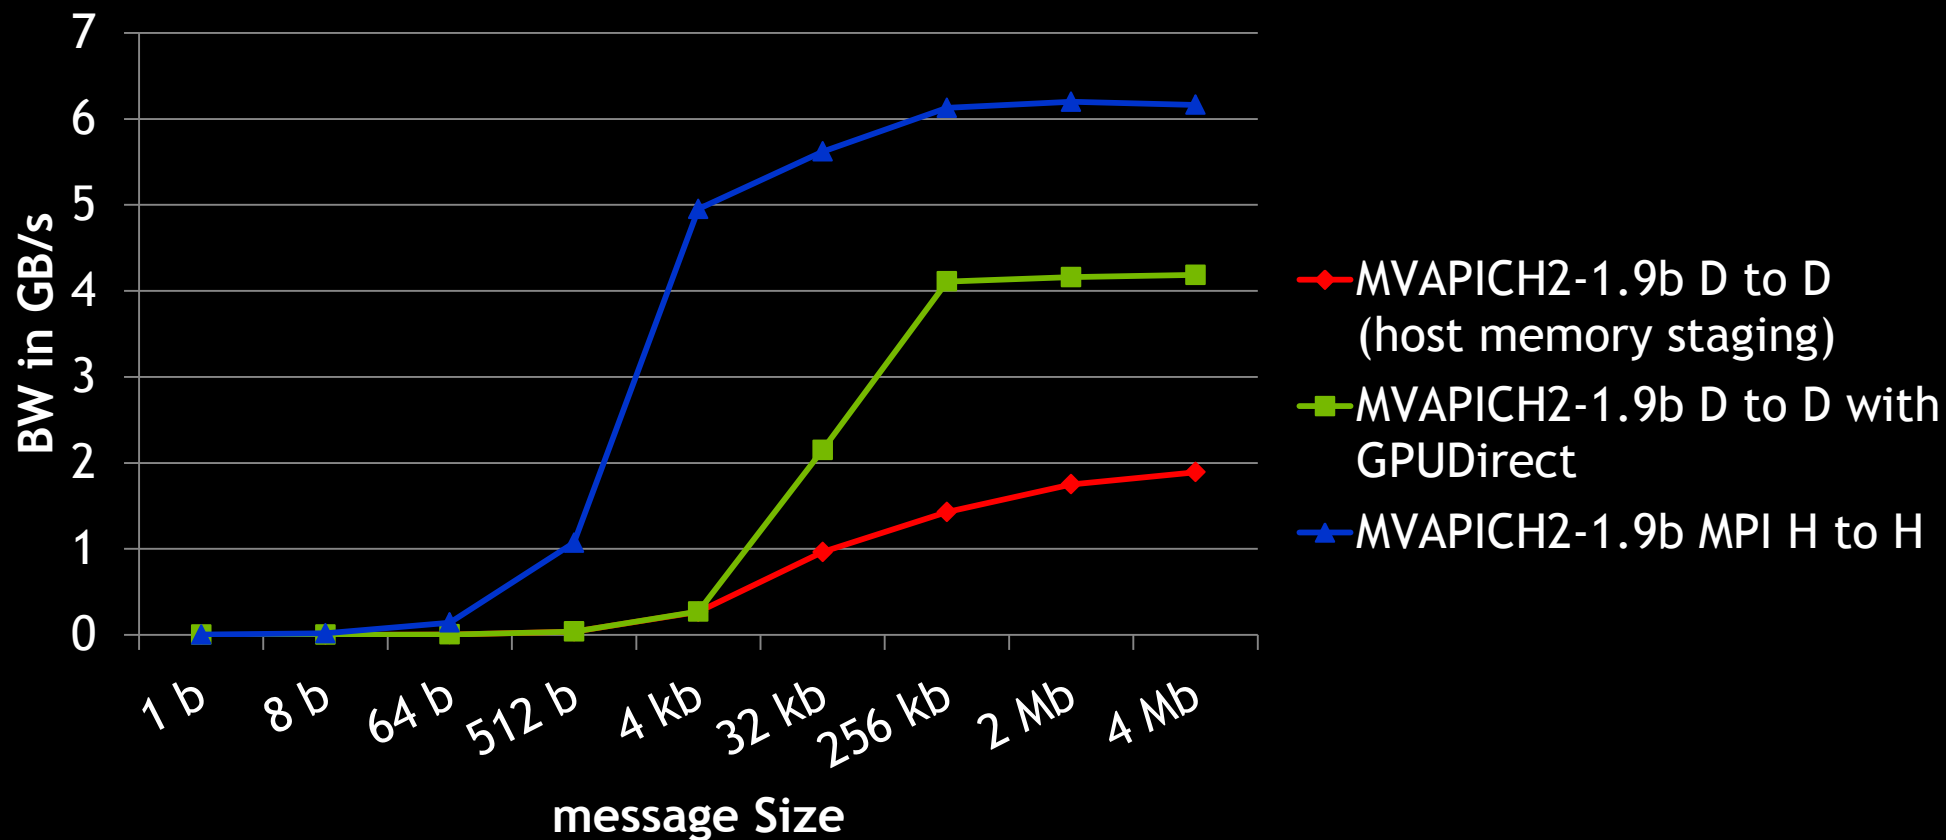
```
MPI_Recv(r_buf_h,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
```


MPI GPU to Remote GPU

Without GPUDirect



Performance Results two Nodes



Latency (1 byte) 19.00 μ s 18.34 μ s 1.11 μ s

Example: Jacobi



- Solves the 2D-Poisson equation on a rectangle

$$\Delta u(x, y) = 0 \quad \forall (x, y) \in \Omega \setminus \delta\Omega$$

- Dirichlet boundary conditions

$$u(x, y) = f(x, y) \in \delta\Omega$$

- 2D domain decomposition with $n \times k$ domains




Example: Jacobi

While not converged

- Do Jacobi step:

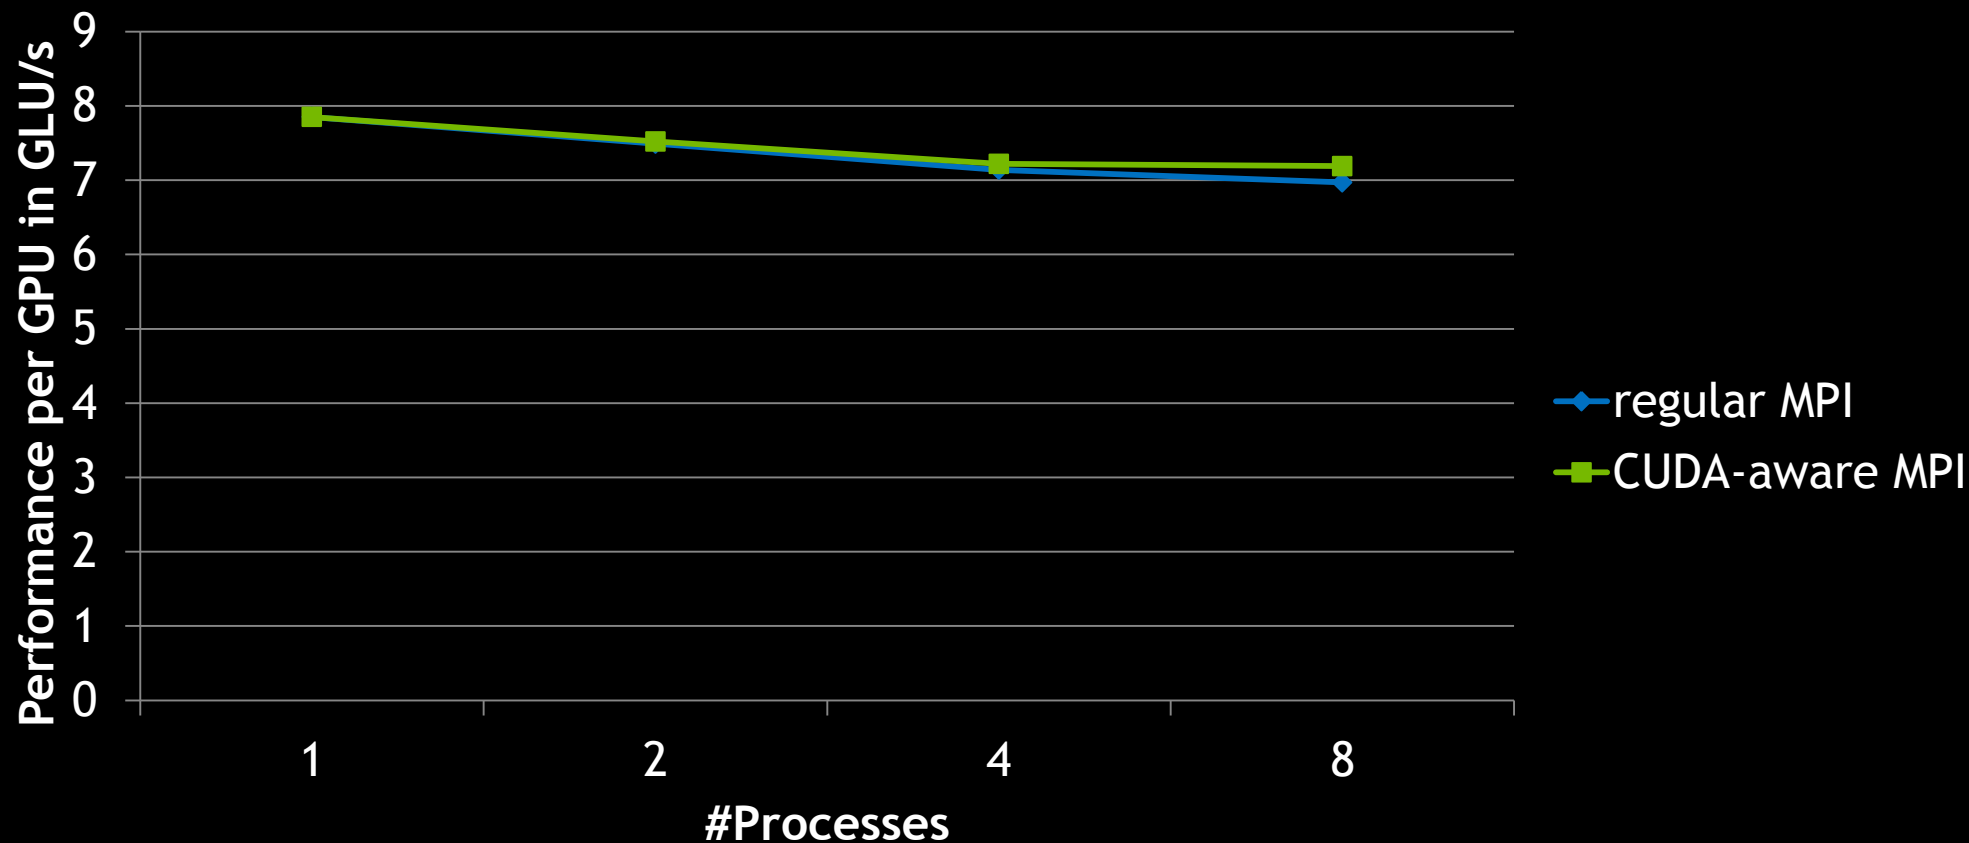
```
for (int i=1; i < n-1; i++) for (int j=1; j < m-1; j++)  
    u_new[i][j] = 0.0f - 0.25f*(u[i-1][j] + u[i+1][j]  
                                + u[i][j-1] + u[i][j+1])
```

- Exchange halo with 2 - 4 neighbours
- Swap `u_new` and `u`
- Next iteration




Jacobi Results (1000 steps)

weak scaling 4k x 4k per process



Jacobi Results (1000 steps)




weak scaling 4k x 4k per process



LBM D2Q37

Lattice Boltzmann Method (LBM)

- D2Q37 Model
- Application developed at U Rome Tore Vergata/INFN, U Ferrara/INFN, TU Eindhoven
- Reproduce dynamics of fluid by simulating virtual particles which collide and propagate, e.g. solve the Rayleigh-Taylor instability
- Simulation of large problems requires double precision and many GPUs

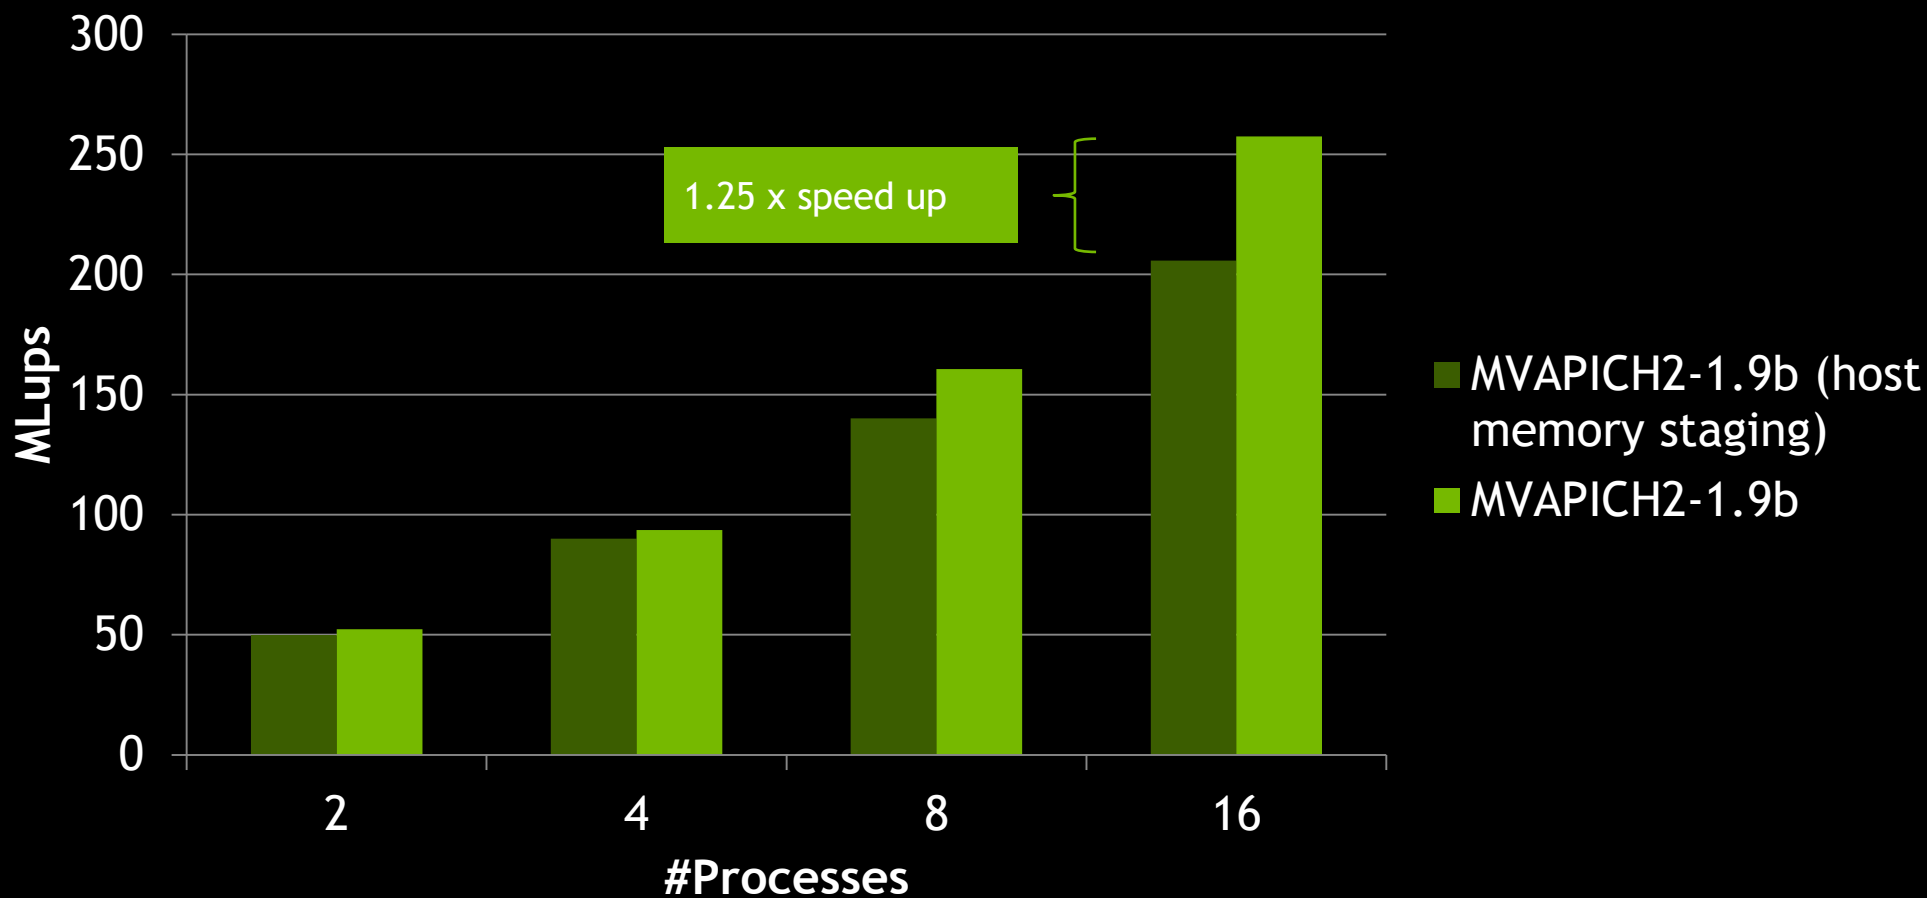


Implementation and Benchmarks:
F. Schifano (U Ferrara)



LBM D2Q37 Results

strong scaling on 8192x1024 cells



CUDA-Aware MPI Implementations

Integrated Support for GPU Computing

- MVAPICH2 1.8/1.9b
 - <http://mvapich.cse.ohio-state.edu/overview/mvapich2/>
- OpenMPI 1.7 (beta)
 - <http://www.open-mpi.org/>
- CRAY MPI (MPT 5.6.2)
- IBM Platform MPI (8.3)

CUDA-Aware Caveats

- `cudaSetDevice` needs to be called before `MPI_Init`
- MPI Environment vars. can be used to set GPU affinity
 - MVAPICH2: `MV2_COMM_WORLD_LOCAL_RANK`
 - OpenMPI: `OMPI_COMM_WORLD_LOCAL_RANK`
- `MV2_USE_CUDA` needs to be set for MVAPICH
- `MPICH_RDMA_ENABLED_CUDA` for MPT on Cray
- `PMPI_GPU_AWARE` for Platform MPI
- Lib needs to be build with CUDA-awareness enabled

CUDA-Aware MPI + OpenACC

- To use CUDA-aware MPI with OpenACC

```
#pragma acc host_data use_device(s_buf)
```

```
MPI_Send(s_buf, size, MPI_CHAR, 1, tag, MPI_COMM_WORLD);
```

- To use MPI with OpenACC

```
#pragma acc update host(s_buf[0:size] )
```

```
MPI_Send(s_buf, size, MPI_CHAR, 1, tag, MPI_COMM_WORLD);
```


Profiling MPI+CUDA applications

- Use nvprof:

```
mpirun -np n nvprof --output-profile out.$MV2_COMM_WORLD_RANK ./app
```

see docs.nvidia.com for details

- Use CUDA-aware tracing libraries like score-p or VampirTrace and tools like Vampir
- Watch the recording of Rolf van de Vaart's Session: S3045 - Tips & Tricks for Getting the Most Out of GPU-accelerated Clusters

Conclusions

- Use CUDA-aware MPI when possible
- Depending on CUDA version, hardware setup, ... a CUDA-aware MPI gives you
 - Ease of programming
 - Pipelined data transfer which automatically provides optimizations when available
 - Overlap CUDA copy and RDMA transfer
 - Utilization of the best GPUDirect technology available
- Examples are available for download at github:


<https://github.com/parallel-forall/code-samples/tree/master/posts/cuda-aware-mpi-example>

Thank you


Examples with source in my Parallel Forall blog posts:

<http://developer.nvidia.com/content/introduction-cuda-aware-mpi>

CUDA Jacobi - Single GPU Performance



Performance Results single Node



Latency (1 byte) 15.87 μ s 19.70 μ s 0.24 μ s